OOREP Project Gap Analysis: Existing vs. LLM-Hermes Vision

Audit date: Saturday, May 30, 2026
Auditor: Hermes Agent (kimi-k2.6)
OOREP path: `~/projects/oorep-local-repertory/`

What EXISTS Right Now

Before listing gaps, here is what is already built and working in your OOREP local repertory project:

Data Layer (COMPLETE)

- **"8,420 rubrics"** from OOREP publicum + Kent OE hierarchy
- **"~2,100 remedies"** with abbreviations, alternative names, IDs
- **"Rubric-to-remedy links"** with classical weight grades (1-3)
- **"Inverted lexical index"** (`rubric_search_index.json`) — word → rubric ID mappings
- **"Vector index"** (`oorep_vector_index.npz`) — 384-dim feature-hashed cosine similarity index

Core API Modules (BUILT)

Module	File	Status	What it does
`HomeopathicRepertory`	`oorep/homeopathic_repertory.py`	✓ Complete	Lexical search, vector search, hybrid search (lexical + vector + overlap), repertorization with classical grade scoring
`ClinicalRubricMapper`	`oorep/clinical_rubric_mapper.py`	✓ Complete	Patient-language normalization, synonym expansion, reviewable rubric suggestion
`RareRemedyTriangulator`	`oorep/rare_remedy_triangulator.py`	✓ Complete	Surfaces small/rare remedies via scarcity metrics
`OOREpVectorSearch`	`oorep/oorep_vector_search.py`	✓ Complete	PHV-1a feature hashing + cosine similarity for rubric retrieval
`RemedyFeedbackStore`	`scripts/remedy_feedback.py`	✓ Complete	SQLite: prescriptions, symptom reports, follow-ups, outcome tracking

Key Capabilities Already Working

- ✓ Natural-language rubric search (lexical)
- ✓ Vector semantic search (offline)
- ✓ Hybrid retrieval (combines both)
- ✓ Repertorization with classical grade-based scoring
- ✓ Patient-language synonym expansion
- ✓ Rare/small remedy surfacing
- ✓ Remedy name and abbreviation lookup
- ✓ Get all rubrics for a remedy, all remedies for a rubric
- ✓ Prescription and outcome tracking via SQLite
- ✓ Data extraction scripts for OOREP dumps
- ✓ pytest test suite

Gap Analysis by Layer

Layer 1: Conversational Repertory Interface

#	Benefit	Status	Gap
1	Natural-language rubric search	✓ EXISTS (lexical + hybrid)	None — working now
2	Graduated rubric exploration	⚠ PARTIAL	Exists in vue data structure, but no parent/child navigation API or CLI
3	Multi-remedy comparison on the fly	✖ MISSING	No API to compare two remedies side-by-side by overlapping rubrics
4	Antidote/complementary lookup	✖ MISSING	No remedy relationships database (complementary, antidotal, inimical)
5	Cross-reference repertory editions	✖ MISSING	Only one OOREP source loaded; no multi-edition comparison
6	Abbreviation decoding	⚠ PARTIAL	`get_remedy_by_abbrev()` exists but doesn't disambiguate collisions robustly

Layer 2: Cross-Session Patient Memory

#	Benefit	Status	Gap
7	Persistent case context	✖ MISSING	`remedy_feedback.py` has SQLite tables but NO Hermes memory integration; you cannot say "what did I prescribe Mrs. J. last month?" in chat
8	Chronic case timelines	✖ MISSING	SQLite has dates but no timeline visualization or query API
9	Patient-specific pattern recognition	✖ MISSING	No cross-case similarity detection
10	Family constellations	✖ MISSING	No family-linking schema in the database
11	Suppression history awareness	✖ MISSING	No dedicated suppression_history table or field

Layer 3: Differential Diagnosis & Selection

#	Benefit	Status	Gap
12	Weighted repertorization dialogue	⚠ PARTIAL	Repertorization exists but not interactive (no live add/remove with updated ranking)
13	Strange-Rare-Peculiar (SRP) flagging	✖ MISSING	No SRP detection logic; all symptoms weighted equally
14	Keynote triangulation	⚠ PARTIAL	"RareRemedyTriangulator" does triangulation but only for rarity, not keynote-specific
15	Elimination analysis	✖ MISSING	No "what symptom rules out X?" logic
16	Potency guidance	✖ MISSING	No potency database or decision tree
17	Acute vs. chronic layer separation	✖ MISSING	No layer tagging in rubrics or symptoms

Layer 4: Materia Medica & Learning

#	Benefit	Status	Gap
18	On-demand proving summaries	✖ MISSING	No materia medica text; only rubric links exist
19	Remedy relationships tutoring	✖ MISSING	No relationships database
20	Source material tracing	✖ MISSING	No provingAuthor attribution in data schema
21	Comparative materia medica	✖ MISSING	No structured materia medica to compare
22	Kingdom/family/group analysis	✖ MISSING	No kingdom/taxonomy classification in remedies.json
23	Clinical tip extraction	✖ MISSING	No NLP mining of outcome notes for success patterns

Layer 5: Pattern Discovery & Research

#	Benefit	Status	Gap
24	Rubric co-occurrence mining	✖ MISSING	No co-occurrence matrix or frequent-itemset analysis
25	Severity-weighted trending	✖ MISSING	No time-series analysis of rubric importance
26	Patient-cohort analysis	✖ MISSING	SQL queries exist but no analytics module
27	Phantom rubric detection	✖ MISSING	No rubric differentiation analysis
28	Botanical-repertory cross-mapping	✖ MISSING	Botanical repertory lives elsewhere; no cross-linking schema
29	Genomic-modality hypothesis	✖ MISSING	No SNP-metabolic data linked to remedy outcomes

Layer 6: Automated Documentation & SOAP

#	Benefit	Status	Gap
30	Voice-to-rubric intake	✖ MISSING	Blue Snowball exists, but no STT-to-rubric pipeline in OOREP
31	SOAP auto-assembly	✖ MISSING	No SOAP generation from conversational notes
32	Prescription audit trails	⚠ PARTIAL	SQLite schema has prescription logs but no immutable audit trail with rubric rationale
33	Follow-up cron scheduling	✖ MISSING	`remedy_feedback.py` tracks `next_followup` but no cron integration
34	Letter generation	✖ MISSING	No document generation module

Layer 7: Multi-Agent & Delegation Workflows

#	Benefit	Status	Gap
35	Rubrics + materia medica + strategy agents	✖ MISSING	No subagent delegation within OOREP
36	Literature-review agent	✖ MISSING	No PubMed/homeopathic journal monitoring
37	Case-supervision agent	✖ MISSING	No "second opinion" re-repertorization agent
38	Student-training agent	✖ MISSING	No simulated patient generation

Layer 8: Data Engineering & Repertory Maintenance

#	Benefit	Status	Gap
39	Automated remedy-picture freshness	✖ MISSING	No update pipeline for new provings
40	Rubric-gap analysis	✖ MISSING	No analysis of symptoms that map poorly to rubrics
41	Custom rubric creation	✖ MISSING	SQLite has `my_private_rubrics` table
42	Vector index rebalancing	⚠ PARTIAL	`build_index()` exists but no automation (cron or file-watcher)
43	Backup & sync verification	✖ MISSING	No GitHub push or integrity verification

Layer 10: Safety, Privacy & Clinical Guardrails

#	Benefit	Status	Gap
44	Materia medica flashcards	✖ MISSING	No flashcard generation or spaced repetition
45	Clinical vignette quizzes	✖ MISSING	No quiz generation engine
46	Kent vs. Boenninghausen comparison	✖ MISSING	Only one method implemented (classical grade-sum)
47	Remedy personality storytelling	✖ MISSING	`personality_engine.py` exists separately but not linked to OOREP
48	Grand rounds synthesis	✖ MISSING	No multi-case aggregation for teaching

Meta-Capabilities

#	Benefit	Status	Gap
49	PHI-minimizing mode	⚠ PARTIAL	You have preference, but no automated PHI scrubbing in OOREP code
50	Practitioner override enforcement	⚠ PARTIAL	`prescriber_ack` field exists but no approval gate in API
51	Red-flag symptom detection	✖ MISSING	No red-flag keyword list or referral trigger logic
52	Contraindicated remedy alerts	✖ MISSING	No remedy-family-reaction history tracking
53	Audit logging for licensure	⚠ PARTIAL	SQLite timestamps exist but no audit-grade immutability
54	Offline-resilience	✓ EXISTS	All data local; no cloud dependency for core functionality

---

Summary Scorecard

#	Benefit	Status	Gap
55	Skill accumulation	✖ MISSING	Hermes saves skills, but no OOREP-specific workflow skills auto-generated
56	Personality-aware reasoning	⚠ PARTIAL	`personality_engine.py` exists elsewhere, not linked to repertory personality
57	Model-routing intelligence	✖ MISSING	No Jetson vs. cloud routing logic in OOREP

Coverage: 3.4% complete, 17.2% partial, 79.3% missing

Recommended Build-Out Priority List

Phase 1: Core Integration (Weeks 1-2)

These plug directly into your existing API and unlock immediate clinical value.

- **"OOREP Hermes Skill Bridge"** — Create a Hermes skill that exposes all `HomeopathicRepertory` methods via natural-language commands. Priority: highest. Unlocks benefits #1, #6, #12 within chat.
- **"Patient Case Manager"** — Extend `RemedyFeedbackStore` with Hermes memory integration so you can query cases by pseudonym, timeline, and outcome across sessions. Covers benefits #7, #8, #9, #11.
- **"Multi-Remedy Comparison API"** — Add `compare_remedies(abbrevA, abbrevB)` to `homeopathic_repertory.py` returning overlap and divergence tables. Covers benefit #3.
- **"Practitioner Approval Gate"** — Enforce `prescriber_ack` with a Hermes `clarify()` prompt before any remedy recommendation is recorded. Covers #50.

Phase 2: Intelligence Layer (Weeks 4-8)

These add analytical depth to existing search and tracking.

- **"SRP Detector Module"** — Build a keyword/modality list of SRP markers and weight them differently in `ClinicalRubricMapper`. Covers #13.
- **"Rubric Co-occurrence Engine"** — Build a frequency matrix of remedy pairs across rubrics; mine for polycrest patterns. Covers #24.
- **"Phantom Rubric Analyzer"** — Compute Giri coefficient of remedies per rubric to flag low-differentiation rubrics. Covers #27.
- **"Patient Cohort Queries"** — SQL views + Python analytics for "most common follow-up remedy for Pulsatilla." Covers #26.
- **"Custom Private Rubrics"** — Add `private_rubrics` table to SQLite with user attribution; integrate into repertorization. Covers #41.

Phase 3: Knowledge Expansion (Weeks 9–14)

These require new data sources.

- **"Remedy Relationships Database"** — Create a JSON/CSV of complementary, antidotal, inimical, antidote relationships from classical materia medica. Covers #4, #19, #20, #21.
- **"Materia Medica Proving Summaries"** — Import or scrape materia medica text (Kent, Boericke, etc.), link to remedy IDs, add full-text search. Covers #18, #21, #22.
- **"Kingdom/Taxonomy Classification"** — Add kingdom, family, group, chemical-column tags to `remedies.json`. Covers #22.
- **"Botanical Cross-Mapping"** — Link OOREP remedy IDs to WHO Monograph IDs and your botanical repertory. Covers #28.
- **"Red-Flag Symptom Detection"** — Maintain a keyword list for symptoms requiring allopathic referral; gate repertorization results when detected. Covers #51.

Phase 4: Workflow Automation (Weeks 15-20)

These are the polish layers for daily practice.

- **"SOAP Auto-Assembly"** — From conversational case notes, generate SOAP draft using LLM + rubric tagging. Covers #31.
- **"Follow-Up Cron Jobs"** — Schedule Hermes cron notifications for prescription follow-ups using `next_followup` field. Covers #33.
- **"Letter Generation"** — Template-based referral letter generation with homeopathic rationale incorporation. Covers #34.
- **"Voice-to-Rubric Pipeline"** — Connect your Blue Snowball STT → clinical rubric mapper → rubric suggestions. Covers #30.
- **"Vector Index Auto-Rebuild"** — File watcher or cron to rebuild the vector index when `rubrics.json` updates. Covers #42.
- **"GitHub Backup Cron"** — Encrypted snapshots of SQLite + JSON data to private repo. Covers #43.

Phase 5: Advanced / Teaching (Beyond 20 weeks)

- **"Subagent Delegation Architecture"** — Spawn delegated agents for rubrics research, materia medica lookup, strategy. Covers #35, #36, #37.
- **"Clinical Vignette Quiz Generator"** — Build case scenarios from your outcome database for self-testing. Covers #45.
- **"Personality Engine Integration"** — Link OOREP remedy IDs to your 30-remedy personality system for narrative teaching. Covers #47, #56.
- **"Genomic-Modality Hypothesis Engine"** — SNP data → remedy outcome correlation mining. Covers #29.
- **"Grand Rounds Synthesizer"** — Aggregate anonymized cases into composite teaching narratives. Covers #48.

Immediate Next Step

If you want to start building today, the **single highest-leverage deliverable** is creating a Hermes skill that wraps the existing `HomeopathicRepertory` API. That skill alone would immediately unlock conversational repertorization, rubric search, remedy lookup, and rare remedy triangulation from within any chat session — turning your OOREP project from a Python library into an interactive clinical tool.

Want me to scaffold that skill now?

Layer	Total Benefits	EXISTS	PARTIAL	MISSING
1. Conversational Repertory	6	1	2	3
2. Patient Memory	5	0	0	5
3. Differential Diagnosis	6	0	3	3
4. Materia Medica	6	0	0	6
5. Pattern Discovery	6	0	0	6
6. Documentation/SOAP	5	0	1	4
7. Multi-Agent Workflows	4	0	0	4
8. Data Engineering	5	0	1	4
9. Teaching/Community	5	0	0	5
10. Safety/Privacy	6	1	2	3
Meta-Capabilities	5	0	1	2
TOTALS	**59**	**1**	**10**	**48**